

AtliQo BANK Credit Card Launch

PHASE 1 - Figure out the Target Market

- Distributions : Normal, Skewness
- Data Cleaning : Handling Null Values
- EDA : Pandas, Seaborn, Matplotlib
- EDA : Measures of Central Tendency
- EDA : Measures of Dispersion
- Outlier Treatment : IQR, StdDev, Mode
- Data Visualization : Histogram, Countplot

```
import pandas as pd
import numpy as np
import seaborn as sns
import matplotlib.pyplot as plt
import warnings
warnings.filterwarnings("ignore")

df_cust = pd.read_csv('E:/Study/Data Analysis/DS - BC 2025/2 - Math &
Statistics/5 - PROJECT/PHASE 1/datasets/customers.csv')
df_cp = pd.read_csv('E:/Study/Data Analysis/DS - BC 2025/2 - Math &
Statistics/5 - PROJECT/PHASE 1/datasets/credit_profiles.csv')
df_trans = pd.read_csv('E:/Study/Data Analysis/DS - BC 2025/2 - Math &
Statistics/5 - PROJECT/PHASE 1/datasets/transactions.csv')
```

```
df_cust.head(5)
```

	cust_id	name	gender	age	location	occupation \
0	1	Manya Acharya	Female	2	City	Business Owner
1	2	Anjali Pandey	Female	47	City	Consultant
2	3	Aaryan Chauhan	Male	21	City	Freelancer
3	4	Rudra Bali	Male	24	Rural	Freelancer
4	5	Advait Malik	Male	48	City	Consultant

	annual_income	marital_status
0	358211.0	Married
1	65172.0	Single
2	22378.0	Married
3	33563.0	Married
4	39406.0	Married

```
df_cp.head(5)
```

	cust_id	credit_score	credit_utilisation	outstanding_debt \
0	1	749	0.585171	19571.0
1	2	587	0.107928	161644.0
2	3	544	0.854807	513.0
3	4	504	0.336938	224.0
4	5	708	0.586151	18090.0

```

    credit_inquiries_last_6_months  credit_limit
0                                0.0         40000.0
1                                2.0         1250.0
2                                4.0         1000.0
3                                2.0         1000.0
4                                2.0         40000.0

df_trans.head(5)

   tran_id  cust_id  tran_date  tran_amount  platform
product_category \
0         1      705  2023-01-01          63  Flipkart
Electronics
1         2      385  2023-01-01          99  Alibaba  Fashion &
Apparel
2         3      924  2023-01-01         471  Shopify
Sports
3         4      797  2023-01-01          33  Shopify  Fashion &
Apparel
4         5      482  2023-01-01          68  Amazon  Fashion &
Apparel

   payment_type
0      Phonepe
1  Credit Card
2      Phonepe
3        Gpay
4  Net Banking

```

SQL Connection

```

import mysql.connector

conn = mysql.connector.connect(
    host='localhost',
    user='root',
    passwd='root',
    database='e_master_card'
)

df_cust = pd.read_sql("SELECT * FROM customers", conn)
df_cust.head(3)

```

	cust_id	name	gender	age	location	occupation \
0	1	Manya Acharya	Female	2	City	Business Owner
1	2	Anjali Pandey	Female	47	City	Consultant
2	3	Aaryan Chauhan	Male	21	City	Freelancer

```

    annual_income marital_status
0         358211      Married
1         65172      Single
2         22378      Married

df_cp = pd.read_sql("SELECT * FROM credit_profiles", conn)
df_cp.head(3)

   cust_id  credit_score  credit_utilisation  outstanding_debt \
0         1           749           0.585171           19571.0
1         2           587           0.107928           161644.0
2         3           544           0.854807             513.0

   credit_inquiries_last_6_months  credit_limit
0                             0.0         40000.0
1                             2.0          1250.0
2                             4.0          1000.0

df_trans = pd.read_sql("SELECT * FROM transactions", conn)
df_trans.head(3)

   tran_id  cust_id  tran_date  tran_amount  platform
product_category \
0         1       705  2023-01-01           63  Flipkart
Electronics
1         2       385  2023-01-01           99  Alibaba  Fashion &
Apparel
2         3       924  2023-01-01          471  Shopify
Sports

   payment_type
0      Phonepe
1  Credit Card
2      Phonepe

# when you are done importing the data, close the connection
conn.close()

```

Overview of Data

```

print("Customers data - ",df_cust.shape)
print("Credit Score data - ",df_cp.shape)
print("Transactions data - ",df_trans.shape)

Customers data - (1000, 8)
Credit Score data - (1004, 6)
Transactions data - (500000, 7)

```

```
df_cust.describe()
```

	cust_id	age	annual_income
count	1000.000000	1000.000000	1000.000000
mean	500.500000	36.405000	132439.799000
std	288.819436	15.666155	113706.313793
min	1.000000	1.000000	0.000000
25%	250.750000	26.000000	42229.750000
50%	500.500000	32.000000	107275.000000
75%	750.250000	46.000000	189687.500000
max	1000.000000	135.000000	449346.000000

```
df_cp.describe()
```

	cust_id	credit_score	credit_utilisation	outstanding_debt
count	1004.000000	1004.000000	1000.000000	1000.000000
mean	500.850598	588.655378	0.498950	9683.597000
std	288.315670	152.575244	0.233139	25255.893671
min	1.000000	300.000000	0.103761	33.000000
25%	251.750000	459.000000	0.293917	221.000000
50%	502.500000	601.000000	0.487422	550.000000
75%	749.250000	737.250000	0.697829	11819.500000
max	1000.000000	799.000000	0.899648	209901.000000

	credit_inquiries_last_6_months	credit_limit
count	1000.000000	935.000000
mean	1.955000	19235.561497
std	1.414559	24489.997195
min	0.000000	500.000000
25%	1.000000	750.000000
50%	2.000000	1250.000000
75%	3.000000	40000.000000
max	4.000000	60000.000000

```
df_trans.describe()
```

	tran_id	cust_id	tran_amount
count	500000.000000	500000.000000	500000.000000
mean	250000.500000	501.400428	3225.20733
std	144337.711635	288.641924	13098.74276
min	1.000000	1.000000	0.000000
25%	125000.750000	252.000000	64.000000

50%	250000.500000	502.000000	141.000000
75%	375000.250000	752.000000	397.000000
max	500000.000000	1000.000000	69999.000000

Data Discrepancy

- For a bank customer, the min. age cannot be 1, and max. age cannot be 135
- The min. annual income for a bank customer, cannot be 0


Therefore, the above discrepancies will be taken as **OUTLIERS**.

Treating Null Values

```
df_cust.isnull().sum()
```

```
cust_id      0
name         0
gender       0
age          0
location     0
occupation   0
annual_income 0
marital_status 0
dtype: int64
```

```
plt.figure(figsize=(5, 5))
sns.histplot(df_cust['annual_income'], kde=True, color='green',
label='Data')
plt.title('Histogram of annual_income')
plt.show()
```



We have following observations from the above,

1. **Age:** min = 1, max = 135
2. **Annual Income:** min = 2, max = 447 k

Age column has outliers. Annual income also seem to have outliers in terms of minimum value because business suggested that minimum income should be atleast 100

```
df_cust.annual_income.describe()
```

```
count      1000.000000
mean      132439.799000
std       113706.313793
min         0.000000
25%       42229.750000
50%      107275.000000
75%      189687.500000
max      449346.000000
Name: annual_income, dtype: float64
```

```
df_cp.isnull().sum()
```

```

cust_id          0
credit_score      0
credit_utilisation 4
outstanding_debt  4
credit_inquiries_last_6_months 4
credit_limit      69
dtype: int64

```

Outlier Detection: Annual income

Let us use standard deviation to detect outliers. Common practice is to treat anything that ± 3 std dev as an outlier

```

df_cust['annual_income'].mean(), df_cust['annual_income'].std()
(132439.799, 113706.31379289791)

lower_sd = df_cust['annual_income'].mean() -
3*df_cust['annual_income'].std()
upper_sd = df_cust['annual_income'].mean() +
3*df_cust['annual_income'].std()

lower_sd, upper_sd
(-208679.14237869374, 473558.74037869374)

df_cust[df_cust['annual_income']>upper_sd]

Empty DataFrame
Columns: [cust_id, name, gender, age, location, occupation,
annual_income, marital_status]
Index: []

```

We are seeing two outliers as per our statistical criteria of ± 3 std dev.

But we don't always assume these as outliers all the time. We have to use business knowledge and our sense of judgement. Here after discussing with the business we concluded that having this type of higher income for business owners is usual and we will keep these data points as is to stay close to the reality while doing our analysis.

On the lower end however, we see minimum income as 2. Our business manager has told us that the income should be at least 100. We can use this as our criteria to find out the outliers on the lower end. These outliers could have occurred due to a data error.

```

df_cust[df_cust.annual_income < 100]

   cust_id  name  gender  age  location
occupation \

```

14	15	Sanjana Malik	Female	25	Rural	
Artist						
31	32	Veer Mistry	Male	50	City	Business
Owner						
82	83	Reyansh Mukherjee	Male	27	City	
Freelancer						
97	98	Virat Puri	Male	47	Suburb	Business
Owner						
102	103	Aarav Shah	Male	32	City	Data
Scientist						
155	156	Kiaan Saxena	Male	24	City	Fullstack
Developer						
170	171	Advait Verma	Male	52	City	Business
Owner						
186	187	Samar Sardar	Male	53	City	
Consultant						
192	193	Ishan Joshi	Male	37	Suburb	Data
Scientist						
227	228	Advait Mukherjee	Male	48	City	Business
Owner						
232	233	Aditya Goel	Male	26	City	
Freelancer						
240	241	Aaryan Bose	Male	24	Suburb	
Freelancer						
262	263	Vivaan Tandon	Male	53	Suburb	Business
Owner						
272	273	Kunal Sahani	Male	50	Suburb	Business
Owner						
275	276	Ananya Bali	Female	47	City	
Consultant						
312	313	Ritvik Gupta	Male	50	City	
Consultant						
315	316	Amara Jha	Female	25	City	Data
Scientist						
316	317	Yuvraj Saxena	Male	47	City	
Consultant						
333	334	Avani Khanna	Female	29	City	Data
Scientist						
340	341	Priya Sinha	Female	33	Rural	Fullstack
Developer						
402	403	Arnav Singh	Male	60	City	Business
Owner						
404	405	Arnav Banerjee	Male	26	City	Data
Scientist						
409	410	Kiaan Jain	Male	45	Rural	
Consultant						
440	441	Rudra Bose	Male	36	Suburb	Data
Scientist						
446	447	Aahan Gambhir	Male	60	City	Business

Owner							
449	450	Anika Rathod	Female	24	Suburb	Fullstack	
Developer							
461	462	Kunal Nair	Male	33	City	Data	
Scientist							
474	475	Neha Verma	Female	28	City	Data	
Scientist							
502	503	Samar Dewan	Male	38	Suburb	Data	
Scientist							
508	509	Advait Das	Male	55	City	Business	
Owner							
516	517	Rehan Kulkarni	Male	29	Rural	Fullstack	
Developer							
530	531	Aarya Ver	Male	32	City	Business	
Owner							
536	537	Ritvik Patil	Male	33	City	Data	
Scientist							
543	544	Advait Batra	Male	54	City		
Consultant							
592	593	Priya Gandhi	Female	32	City	Business	
Owner							
599	600	Ishan Goswami	Female	38	City		
Consultant							
603	604	Kunal Malhotra	Male	25	Suburb	Fullstack	
Developer							
608	609	Kriti Lalwani	Female	25	City	Data	
Scientist							
633	634	Rudra Mehtani	Male	26	City	Data	
Scientist							
634	635	Anaya Dutta	Female	21	City		
Freelancer							
644	645	Dhruv Das	Male	64	City	Business	
Owner							
648	649	Kunal Rathore	Male	41	City		
Consultant							
650	651	Gauri Mittal	Female	47	Rural		
Consultant							
664	665	Ayush Khanna	Male	32	Rural	Fullstack	
Developer							
681	682	Arya Jaiswal	Male	37	Suburb	Data	
Scientist							
686	687	Vihaan Jaiswal	Male	40	City	Business	
Owner							
688	689	Dhruv Dewan	Male	26	City		
Artist							
693	694	Aditi Mehrotra	Female	37	Suburb	Data	
Scientist							
694	695	Rohan Mehta	Male	28	City	Data	
Scientist							

696	697	Ishan Negi	Male	47	City	
Consultant						
744	745	Swara Kaul	Female	39	City	Data
Scientist						
784	785	Rohan Jain	Male	27	City	Data
Scientist						
788	789	Vihaan Singhal	Male	20	City	Fullstack
Developer						
791	792	Sara Mhatre	Female	38	City	Data
Scientist						
817	818	Akshay Mehrotra	Male	47	City	
Consultant						
932	933	Avinash Tiwari	Male	35	City	Data
Scientist						
955	956	Aahan Gandhi	Male	39	Suburb	Business
Owner						
956	957	Priya Malik	Female	24	City	
Artist						
995	996	Manya Vasudeva	Female	26	City	
Freelancer						
998	999	Amara Rathore	Female	47	City	Business
Owner						

	annual_income	marital_status
14	0	Married
31	50	Married
82	0	Single
97	0	Married
102	0	Married
155	0	Married
170	0	Single
186	0	Single
192	0	Married
227	0	Married
232	0	Married
240	0	Married
262	50	Married
272	0	Married
275	0	Single
312	0	Married
315	0	Married
316	50	Married
333	50	Married
340	50	Married
402	0	Married
404	0	Single
409	0	Married
440	0	Married
446	0	Married

449	0	Married
461	0	Married
474	0	Single
502	0	Single
508	0	Married
516	0	Single
530	0	Married
536	0	Married
543	2	Married
592	50	Married
599	0	Single
603	0	Married
608	0	Single
633	2	Married
634	0	Married
644	0	Single
648	0	Married
650	0	Married
664	0	Married
681	0	Married
686	2	Married
688	0	Married
693	0	Married
694	0	Married
696	20	Married
744	0	Married
784	0	Single
788	0	Single
791	0	Single
817	0	Single
932	0	Married
955	0	Married
956	0	Married
995	0	Married
998	0	Married

Outlier Treatment: Annual income

Above records (with <100\$ income) are outliers. We have following options to treat them :

1. **Remove them:** After discussion with business, we decided not to remove them as these are valid customers and we want to include them in our analysis
2. **Replace them with mean or median :** Mean is sensitive to outliers. It is better to use median for income values
3. **Replace them with occupation wise median:** Income level may vary based on occupation. For example median income for data scientist can be different from a

median income of a business owner. It is better to use occupation wise median income for replacement

```
occ_wise_inc_median = df_cust.groupby("occupation")
["annual_income"].median()
occ_wise_inc_median
```

```
occupation
Accountant      65265.0
Artist          44915.0
Business Owner  254881.0
Consultant      51175.0
Data Scientist  127889.0
Freelancer      45189.5
Fullstack Developer  74457.0
Name: annual_income, dtype: float64
```

```
occ_wise_inc_median['Artist']
44915.0
```

```
for index, row in df_cust.iterrows():
    if row["annual_income"] < 100:
        occupation = df_cust.at[index, "occupation"]
        df_cust.at[index, "annual_income"] =
occ_wise_inc_median[occupation]
```

```
df_cust[df_cust.annual_income < 100]
```

Empty DataFrame

Columns: [cust_id, name, gender, age, location, occupation, annual_income, marital_status]
Index: []

```
df_cust.loc[[240, 474, 502]]    # Now the customers with income <
$100, is updated with a median occupation income
```

	cust_id	name	gender	age	location	occupation	\
240	241	Aaryan Bose	Male	24	Suburb	Freelancer	
474	475	Neha Verma	Female	28	City	Data Scientist	
502	503	Samar Dewan	Male	38	Suburb	Data Scientist	

	annual_income	marital_status
240	45189.5	Married
474	127889.0	Single
502	127889.0	Single


- The customers having income < \$100 can be seen in **cell[21]**
- The 'median' of the occ. wise income can be seen in **cell[27]**

Data Visualization - Annual Income

```
avg_income_per_occ = df_cust.groupby("occupation")
["annual_income"].mean().round(2)
avg_income_per_occ

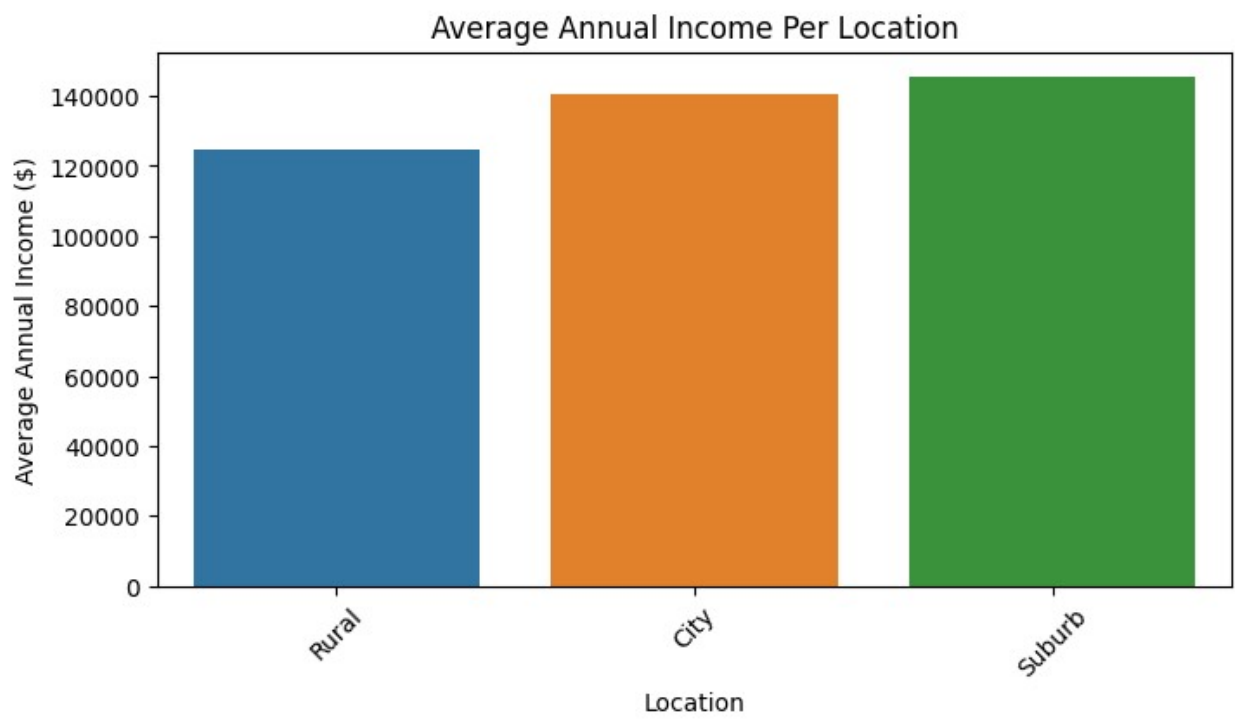

occupation
Accountant      64123.56
Artist          45239.84
Business Owner  268119.83
Consultant      59927.26
Data Scientist  136208.60
Freelancer      76293.09
Fullstack Developer  78618.39
Name: annual_income, dtype: float64

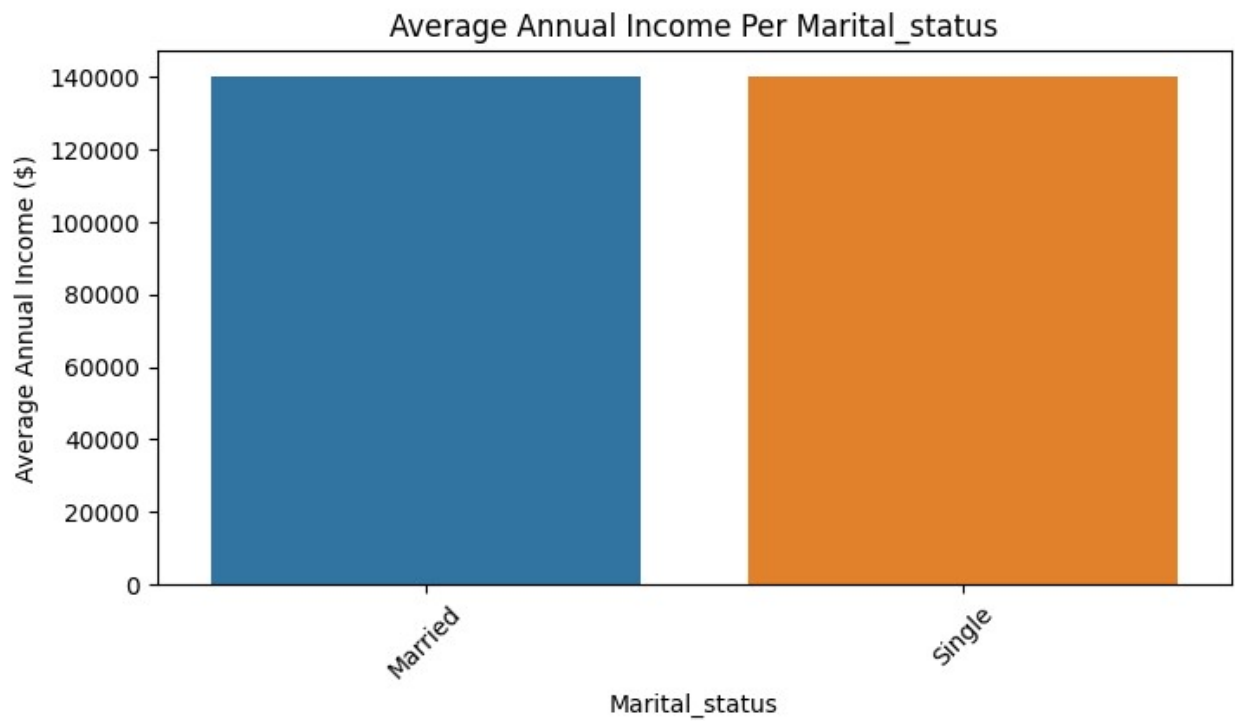
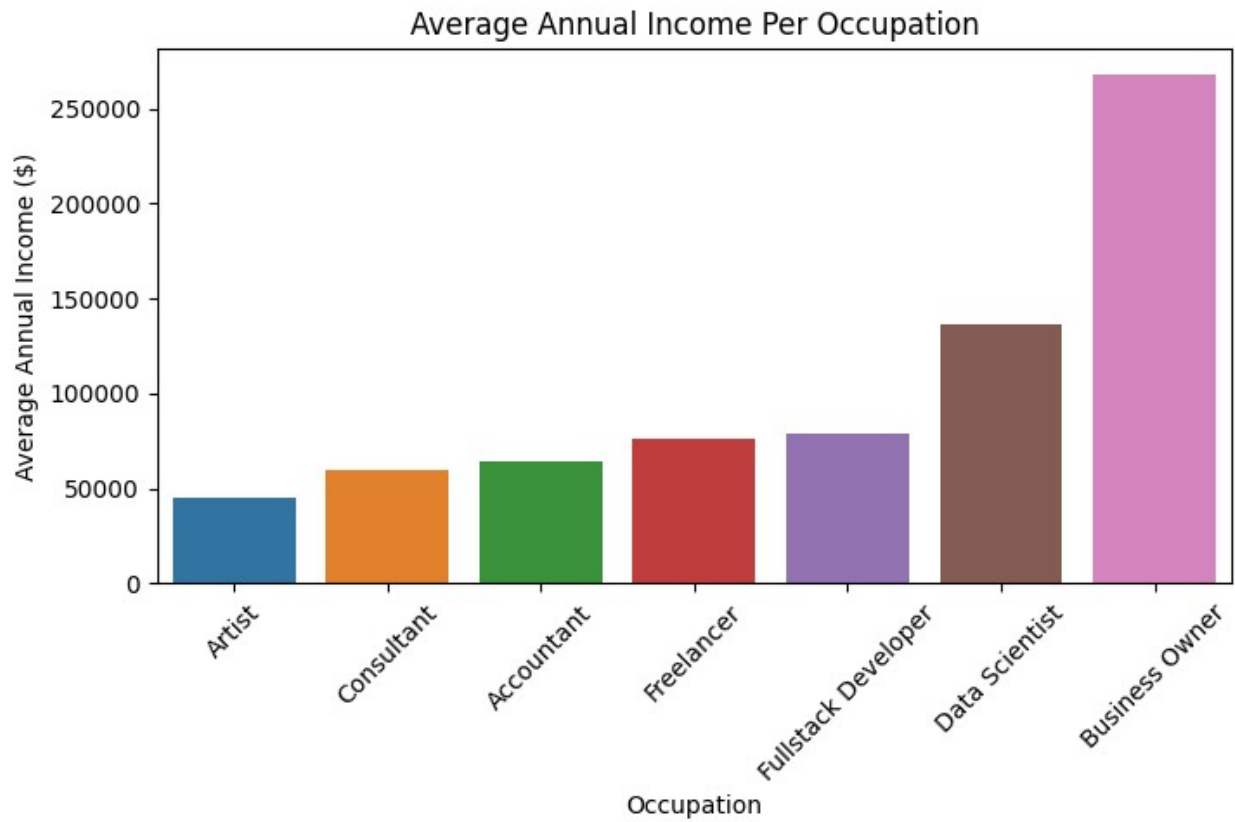
plt.figure(figsize = (8,4))
sns.barplot(x = avg_income_per_occ.index, y =
avg_income_per_occ.values, palette = 'tab10')
plt.xticks(rotation = 45)
plt.title('Average Annual Income Per Occupation')
plt.xlabel('Occupation')
plt.ylabel('Average Annual Income ($)')
plt.show()
```



```
# List of categorical columns
categorical_columns = ['gender', 'location', 'occupation',
                       'marital_status']

# Loop through each categorical column and plot a bar chart of average
# annual income
for col in categorical_columns:
    plt.figure(figsize = (8, 4))
    avg_income_per_group = df_cust.groupby(col)
    ['annual_income'].mean().sort_values()
    sns.barplot(x = avg_income_per_group.index, y =
avg_income_per_group.values, palette = 'tab10')
    plt.xticks(rotation=45)
    plt.title(f'Average Annual Income Per {col.capitalize()}')
    plt.xlabel(col.capitalize())
    plt.ylabel('Average Annual Income ($)')
    plt.show()
```





Analysis of Age Column

```
df_cust.age.isnull().sum()
```

```
0
```

```
df_cust.describe()
```

	cust_id	age	annual_income
count	1000.000000	1000.000000	1000.000000
mean	500.500000	36.405000	140137.395500
std	288.819436	15.666155	110450.464107
min	1.000000	1.000000	5175.000000
25%	250.750000	26.000000	49620.500000
50%	500.500000	32.000000	115328.000000
75%	750.250000	46.000000	195514.250000
max	1000.000000	135.000000	449346.000000

Outlier Treatment: Age

Above we see that min age is 1 and max age is 135. These seem to be outliers. So let's find out age distribution.

```
min_age = df_cust.age.min()
```

```
max_age = df_cust.age.max()
```

```
min_age, max_age
```

```
(1, 135)
```

```
plt.hist(df_cust.age, bins=20, edgecolor='black')
```

```
plt.xlabel("Age")
```

```
plt.ylabel("Count")
```


```
plt.title("Customer Age Distribution")
```

```
plt.axvline(min_age, color="red", label=f"Min Age: {min_age}")
```

```
plt.axvline(max_age, color="green", label=f"Max Age: {max_age}")
```

```
plt.legend()
```

```
plt.show()
```



```
df_cust[(df_cust.age<15)|(df_cust.age>80)] # Client told that age
< 15 and age > 80 are invalid for the data
```

cust_id	name	gender	age	location	occupation
0	Manya Acharya	Female	2	City	Business Owner
41	Aaryan Shah	Male	110	City	Artist
165	Sia Dutta	Female	1	City	Freelancer
174	Rohan Sharma	Male	110	City	Freelancer
222	Arjun Batra	Male	110	Suburb	Freelancer
277	Aarav Tandon	Male	110	City	Consultant
295	Ayush Pandey	Male	1	Rural	Accountant
325	Virat Goel	Male	110	City	Accountant
610	Rehan Verma	Male	135	Rural	Business Owner

692	693	Dhruv Jha	Male	1	City	Business
Owner						
703	704	Aanya Sharma	Female	110	City	
Freelancer						
709	710	Anika Verma	Female	110	City	Data
Scientist						
728	729	Rehan Yadav	Male	135	City	Business
Owner						
832	833	Ridhi Raj	Female	110	City	Fullstack
Developer						
845	846	Rohan Jaiswal	Male	1	City	
Consultant						
855	856	Aanya Taneja	Female	2	City	Fullstack
Developer						
895	896	Krishna Goswami	Male	1	City	
Freelancer						
923	924	Kunal Patel	Male	110	City	
Freelancer						
951	952	Virat Shetty	Male	135	City	Data
Scientist						
991	992	Arya Dube	Male	135	City	Fullstack
Developer						

	annual_income	marital_status
0	358211.0	Married
41	7621.0	Married
165	39721.0	Single
174	23723.0	Married
222	210987.0	Married
277	96522.0	Single
295	55254.0	Married
325	61021.0	Single
610	444776.0	Married
692	83045.0	Married
703	43404.0	Single
709	98417.0	Married
728	382836.0	Married
832	95379.0	Single
845	20838.0	Married
855	30689.0	Married
895	31533.0	Married
923	51629.0	Married
951	49677.0	Married
991	93267.0	Single

```

outliers = df_cust[(df_cust.age<15)|(df_cust.age>80)] # storing
outliers in a separate DF
outliers.shape

(20, 8)

```

Total 20 outliers for age. Now how can we handle these outliers?

Possible options,

- Remove them: This doesn't sound like a good option as we will lose important information
- Replace outlier values with some appropriate value: We can use mean or median for this

```
df_cust.age.median()
```

```
32.0
```

Instead of replacing it with a **median** age for all customers, how about we calculate **median age per occupation**?

```
# Just like we did for the Income part
```

```
median_age_per_occupation = df_cust.groupby('occupation')  
['age'].median()  
median_age_per_occupation
```

```
occupation  
Accountant      31.5  
Artist          26.0  
Business Owner  51.0  
Consultant      46.0  
Data Scientist  32.0  
Freelancer      24.0  
Fullstack Developer  27.5  
Name: age, dtype: float64
```

```
for index, row in outliers.iterrows():  
    if pd.notnull(row['age']):  
        occupation = df_cust.at[index, 'occupation']  
        df_cust.at[index, 'age'] =  
median_age_per_occupation[occupation]
```

```
for index, row in outliers.iterrows():  
    if pd.notnull(row['age']):  
        occupation = df_cust.at[index, 'occupation']  
        df_cust.at[index, 'age'] =  
median_age_per_occupation[occupation]
```

```
df_cust[(df_cust.age<15)|(df_cust.age>80)]
```

```
Empty DataFrame
```

```
Columns: [cust_id, name, gender, age, location, occupation,  
annual_income, marital_status]
```

```
Index: []
```

```
df_cust.age.describe()
count      1000.000000
mean        35.541500
std         12.276634
min         18.000000
25%         26.000000
50%         32.000000
75%         44.250000
max         64.000000
Name: age, dtype: float64
```

We can see above, we don't have any outliers left. Min age is 18 and Max age is 64

Data Visualization - Age and Gender

```
# Defining the bin edges and labels
bin_edges = [17, 25, 48, 65]
bin_labels = ['18-25', '26-48', '49-65']

# Using the cut function to bin and label the age column
pd.cut(df_cust['age'], bins=bin_edges, labels=bin_labels)

0      49-65
1      26-48
2      18-25
3      18-25
4      26-48
...
995    26-48
996    49-65
997    26-48
998    26-48
999    26-48
Name: age, Length: 1000, dtype: category
Categories (3, object): ['18-25' < '26-48' < '49-65']

bin_edges = [17, 25, 48, 65]
bin_labels = ['18-25', '26-48', '49-65']

df_cust['age_group'] = pd.cut(df_cust['age'], bins=bin_edges,
labels=bin_labels)

df_cust['age_group'].value_counts(normalize=True)*100
```


```

age_group
26-48    56.7
18-25    24.6
49-65    18.7
Name: proportion, dtype: float64

# Calculate the count of values in each age group
age_group_counts = df_cust['age_group'].value_counts(normalize=True) *
100

# Plot the pie chart
plt.figure(figsize=(4, 4))
plt.pie(
    age_group_counts,
    labels=age_group_counts.index,
    explode=(0.1,0,0),
    autopct='%1.1f%%',
    shadow=True,
    startangle=140)
plt.axis('equal') # Equal aspect ratio ensures that pie is drawn as a
circle.
plt.title('Distribution of Age Groups')
plt.show()

```



More than 50% of customer base are in the age group of 26 - 48 and ~26% are of age group 18 - 25

Analyze Gender and Location Distribution

```
customer_location_gender = df_cust.groupby(['location',
'gender']).size().unstack(fill_value=0)


# Create a stacked bar chart to visualize the distribution of payment
types for each occupation
customer_location_gender.plot(kind='bar', stacked=True, figsize=(5,
4))

# Add labels and title
plt.xlabel('Location')
plt.ylabel('Count')
plt.title('Customer Distribution by Location and Gender')

# Show the bar chart
plt.legend(title='Payment Type', bbox_to_anchor=(1, 1)) # Add a
legend

# Rotate the x-axis labels for better readability
plt.xticks(rotation=45)

plt.show()
```



Exploring Credit_Score Table - 1

1 - Removing Duplicates

```
df_cp.head()
```

	cust_id	credit_score	credit_utilisation	outstanding_debt	\
0	1	749	0.585171	19571.0	
1	2	587	0.107928	161644.0	
2	3	544	0.854807	513.0	
3	4	504	0.336938	224.0	
4	5	708	0.586151	18090.0	
	credit_inquiries_last_6_months		credit_limit		
0	0.0		40000.0		
1	2.0		1250.0		
2	4.0		1000.0		
3	2.0		1000.0		
4	2.0		40000.0		

```
df_cp.shape
```

```
(1004, 6)
```



```
df_cust.shape
(1000, 9)
df_cp['cust_id'].nunique()
1000
```

There are **4 extra** values in **credit profile**

```
df_cp.duplicated('cust_id')
0      False
1      False
2      False
3      False
4      False
...
999    False
1000   False
1001   False
1002   False
1003   False
Length: 1004, dtype: bool

df_cp[df_cp.duplicated('cust_id', keep=False)]
```

	cust_id	credit_score	credit_utilisation	outstanding_debt	\
516	517	308	NaN	NaN	
517	517	308	0.113860	33.0	
569	569	344	NaN	NaN	
570	569	344	0.112599	37.0	
607	606	734	NaN	NaN	
608	606	734	0.193418	4392.0	
664	662	442	NaN	NaN	
665	662	442	0.856039	266.0	

	credit_inquiries_last_6_months	credit_limit
516	NaN	NaN
517	3.0	500.0
569	NaN	NaN
570	0.0	500.0
607	NaN	NaN
608	1.0	40000.0
664	NaN	NaN
665	2.0	500.0

```
df_cp_upd_1 = df_cp.drop_duplicates(subset='cust_id', keep="last")
df_cp_upd_1.shape
(1000, 6)
```

```
df_cp_upd_1[df_cp_upd_1.duplicated('cust_id', keep=False)]
```

Empty DataFrame

Columns: [cust_id, credit_score, credit_utilisation, outstanding_debt, credit_inquiries_last_6_months, credit_limit]

Index: []

df_cp_upd_1 looks clean now after cleaning duplicates.

Next step would be to see if there are any **null values**

2 - Handling Null Values

```
df_cp_upd_1.isnull().sum()
```

```
cust_id                0
credit_score            0
credit_utilisation      0
outstanding_debt        0
credit_inquiries_last_6_months  0
credit_limit            65
dtype: int64
```

credit_limit has a bunch of null values.

From the business knowledge we know that credit limit depends on credit score of a customer.

We will try to find out if we can figure out a mathematical relationship between credit score and credit limit and use credit score to fill NULL values in credit limit.

```
df_cp_upd_1[df_cp_upd_1.credit_limit.isnull()]
```

	cust_id	credit_score	credit_utilisation	outstanding_debt	\
10	11	679	0.557450	9187.0	
35	36	790	0.112535	4261.0	
37	38	514	0.296971	238.0	
45	46	761	0.596041	24234.0	
64	65	734	0.473715	13631.0	
...	
912	909	479	0.487555	320.0	
931	928	311	0.832244	316.0	
948	945	526	0.272734	227.0	
954	951	513	0.175914	131.0	
957	954	783	0.867421	46451.0	
	credit_inquiries_last_6_months	credit_limit			
10	2.0	NaN			

35	1.0	NaN
37	2.0	NaN
45	2.0	NaN
64	0.0	NaN
..
912	3.0	NaN
931	2.0	NaN
948	1.0	NaN
954	3.0	NaN
957	0.0	NaN

[65 rows x 6 columns]

```
df_cp_upd_1['credit_limit'].unique()
```

```
array([40000., 1250., 1000., 500., 750., nan, 1500., 60000.,
       20000.])
```

```
df_cp_upd_1['credit_limit'].value_counts()
```

```
credit_limit
```

```
500.0      229
```

```
60000.0    186
```

```
40000.0    137
```

```
1500.0     100
```

```
1000.0      90
```

```
750.0       76
```

```
1250.0      75
```

```
20000.0     42
```

```
Name: count, dtype: int64
```

```
# Looking at scatter plot for credit score vs credit_limit again
(after handling outliers)
```

```
# Creating the plot
```

```
plt.figure(figsize=(20, 5))
```

```
plt.scatter(df_cp_upd_1['credit_limit'], df_cp_upd_1['credit_score'],
c='blue', marker='o', label='Data Points')
```

```
# Customizing the plot
```

```
plt.title('Credit Score vs. Credit Limit')
```

```
plt.xlabel('Credit Limit')
```

```
plt.ylabel('Credit Score')
```

```
# Adjusting the y-axis bin interval to 1000
```


```
plt.xticks(range(0, 90001, 5000))
```

```
plt.grid(True)
```

```
# Showing the plot
```

```
plt.legend()
```

```
plt.show()
```



Above, we can see clear relationship between credit score and credit limit.

Where there are levels for example, upto 650 score is getting a very minor credit limit (<1000\$) whereas a score between 650 to 700 is getting around 20000. Score between 700 to 750 is getting around 40K etc.

```
# Defining the bin ranges
bin_ranges = [300, 450, 500, 550, 600, 650, 700, 750, 800]

# Creating labels for the bins
bin_labels = [f'{start}-{end-1}' for start, end in zip(bin_ranges,
bin_ranges[1:])]

# Using pd.cut to assign data to bins
df_cp_upd_1['credit_score_range'] =
pd.cut(df_cp_upd_1['credit_score'], bins=bin_ranges,
labels=bin_labels, include_lowest=True, right=False)

df_cp_upd_1.head()
```

	cust_id	credit_score	credit_utilisation	outstanding_debt \
0	1	749	0.585171	19571.0
1	2	587	0.107928	161644.0
2	3	544	0.854807	513.0
3	4	504	0.336938	224.0
4	5	708	0.586151	18090.0

	credit_inquiries_last_6_months	credit_limit	credit_score_range
0	0.0	40000.0	700-749
1	2.0	1250.0	550-599
2	4.0	1000.0	500-549
3	2.0	1000.0	500-549
4	2.0	40000.0	700-749

We can now see a **new column** called **credit_score_range** which is calculated based on the **credit_score** column

```
df_cp_upd_1[['credit_score', 'credit_score_range',
'credit_limit']].head(3)
```

	credit_score	credit_score_range	credit_limit
0	749	700-749	40000.0
1	587	550-599	1250.0
2	544	500-549	1000.0

```
df_cp_upd_1[df_cp_upd_1['credit_score_range']=="750-799"]
```

	cust_id	credit_score	credit_utilisation	outstanding_debt \
21	22	785	0.897089	36083.0
25	26	758	0.250811	190838.0
26	27	766	0.830908	31344.0
29	30	798	0.222597	7238.0
31	32	768	0.747793	35109.0
...
988	985	770	0.628088	33405.0
993	990	772	0.259958	11937.0
996	993	782	0.477170	20305.0
1000	997	774	0.465462	17139.0
1003	1000	775	0.696050	33956.0

	credit_inquiries_last_6_months	credit_limit	credit_score_range
21	3.0	60000.0	750-799
25	2.0	60000.0	750-799
26	3.0	60000.0	750-799
29	2.0	60000.0	750-799
31	2.0	60000.0	750-799
...
988	2.0	60000.0	750-799
993	2.0	60000.0	750-799
996	2.0	60000.0	750-799
1000	0.0	60000.0	750-799
1003	1.0	60000.0	750-799

```
[213 rows x 7 columns]
```

```
df_cp_upd_1[df_cp_upd_1['credit_score_range']=="300-449"]
```

	cust_id	credit_score	credit_utilisation	outstanding_debt	\
5	6	442	0.705409	246.0	
11	12	429	0.627645	263.0	
15	16	347	0.531660	190.0	
18	19	447	0.795650	292.0	
20	21	381	0.714710	307.0	
...	
981	978	371	0.435307	183.0	
982	979	332	0.150815	65.0	
984	981	327	0.377202	108.0	
989	986	425	0.178470	56.0	
998	995	360	0.594345	242.0	

	credit_inquiries_last_6_months	credit_limit	credit_score_range
5	4.0	500.0	300-449
11	0.0	500.0	300-449
15	0.0	500.0	300-449
18	1.0	500.0	300-449
20	0.0	500.0	300-449
...
981	2.0	500.0	300-449
982	1.0	500.0	300-449
984	3.0	500.0	300-449
989	4.0	500.0	300-449
998	0.0	500.0	300-449

[237 rows x 7 columns]

Above we can see that for credit score range "750-799" the credit limit is 60K whereas for "300-449" it is 500.

We can use MODE function to find out most frequently occurring credit limit for a given score range.

```
mode_df = df_cp_upd_1.groupby('credit_score_range')
['credit_limit'].agg(lambda x: x.mode().iloc[0]).reset_index()
mode_df
```

	credit_score_range	credit_limit
0	300-449	500.0
1	450-499	750.0
2	500-549	1000.0
3	550-599	1250.0
4	600-649	1500.0
5	650-699	20000.0
6	700-749	40000.0
7	750-799	60000.0

```
df_cp_upd_1[df_cp_upd_1.credit_limit.isnull()].sample(3)
```

	cust_id	credit_score	credit_utilisation	outstanding_debt	\
430	431	610	0.741063	628.0	
83	84	733	0.525567	16663.0	
650	648	405	0.231599	63.0	

	credit_inquiries_last_6_months	credit_limit	credit_score_range
430	4.0	NaN	600-649
83	1.0	NaN	700-749
650	0.0	NaN	300-449

Merging the mode values back with the original DataFrame

```
df_cp_upd_2 = pd.merge(df_cp_upd_1, mode_df, on='credit_score_range',
suffixes=('', '_mode'))
df_cp_upd_2.sample(3)
```

	cust_id	credit_score	credit_utilisation	outstanding_debt	\
483	484	708	0.244153	7599.0	
797	798	725	0.766200	23055.0	
341	342	625	0.647972	611.0	

	credit_inquiries_last_6_months	credit_limit	credit_score_range	\
483	0.0	40000.0	700-749	
797	2.0	40000.0	700-749	
341	4.0	1500.0	600-649	

	credit_limit_mode
483	40000.0
797	40000.0
341	1500.0

```
df_cp_upd_2[df_cp_upd_2.credit_limit.isnull()].sample(3)
```

	cust_id	credit_score	credit_utilisation	outstanding_debt	\
849	850	787	0.293520	11195.0	
841	842	490	0.555309	249.0	
114	115	619	0.128910	151.0	

	credit_inquiries_last_6_months	credit_limit	credit_score_range	\
849	3.0	NaN	750-799	
841	1.0	NaN	450-499	
114	1.0	NaN	600-649	

	credit_limit_mode
849	60000.0
841	750.0
114	1500.0

Above, we can simply replace NaN value in credit_limit column with credit_limit_mode value.

This value indicates most frequently occurring credit limit for a given credit_score_range. Hence it can be used as a replacement value.

We will create a new copy of the dataframe so that we have reproducibility and access of the older dataframe in this notebook.

```
df_cp_upd_3 = df_cp_upd_2.copy()
df_cp_upd_3['credit_limit'].fillna(df_cp_upd_3['credit_limit_mode'],
inplace=True)
df_cp_upd_3.shape
(1000, 8)
df_cp_upd_3.isnull().sum()
cust_id          0
credit_score      0
credit_utilisation 0
outstanding_debt  0
credit_inquiries_last_6_months 0
credit_limit      0
credit_score_range 0
credit_limit_mode 0
dtype: int64
```

You can now see **ZERO outliers** in credit_limit column which means we successfully got rid of all NULL values.

```
df_cp_upd_3[df_cp_upd_3.cust_id==431]
```

	cust_id	credit_score	credit_utilisation	outstanding_debt	\
430	431	610	0.741063	628.0	

	credit_inquiries_last_6_months	credit_limit	credit_score_range	\
430	4.0	1500.0	600-649	

	credit_limit_mode
430	1500.0

Previously customer id **431** had null value in credit_limit. Now it has a valid value.

3 - Handling Outliers

```
# outstanding_debt
```

```
df_cp_upd_3.describe()
```

	cust_id	credit_score	credit_utilisation	outstanding_debt
\				
count	1000.000000	1000.000000	1000.000000	1000.000000
mean	500.500000	589.182000	0.498950	9683.597000
std	288.819436	152.284929	0.233139	25255.893671
min	1.000000	300.000000	0.103761	33.000000
25%	250.750000	460.000000	0.293917	221.000000
50%	500.500000	601.500000	0.487422	550.000000
75%	750.250000	738.000000	0.697829	11819.500000
max	1000.000000	799.000000	0.899648	209901.000000


	credit_inquiries_last_6_months	credit_limit	credit_limit_mode
count	1000.000000	1000.000000	1000.000000
mean	1.955000	19733.75000	19912.500000
std	1.414559	24717.43818	24840.914633
min	0.000000	500.00000	500.000000
25%	1.000000	750.00000	750.000000
50%	2.000000	1500.00000	1500.000000
75%	3.000000	40000.00000	40000.000000
max	4.000000	60000.00000	60000.000000

When we observe min and max for various columns, we realize that **outstanding_debt's max > the max of credit_limit**.

Based on the business understanding, we know that the maximum debt that a customer can have is equal to credit limit. They would not be allowed to spend more than their credit limit.

```
# Checking Outliers
```

```
plt.figure(figsize=(5, 5))  
sns.boxplot(x = df_cp_upd_3['outstanding_debt'])  
plt.title('Box plot for outstanding debt')  
Text(0.5, 1.0, 'Box plot for outstanding debt')
```



Instead of using any statistical approach (such as standard deviation or IQR), here too we will use a business knowledge. We will mark any outstanding debt that is greater than credit limit as an outlier

And, we will replace these outliers with credit_limit. We can assume that there was some data processing error due to we got these high numbers and it is ok to replace them with a credit_limit

```
df_cp_upd_3[df_cp_upd_3.outstanding_debt>df_cp_upd_3.credit_limit]
```

	cust_id	credit_score	credit_utilisation	outstanding_debt \
1	2	587	0.107928	161644.0
19	20	647	0.439132	205014.0
25	26	758	0.250811	190838.0

38	39	734	0.573023	122758.0
93	94	737	0.739948	137058.0
204	205	303	0.364360	187849.0
271	272	703	0.446886	154568.0
301	302	722	0.608076	122402.0
330	331	799	0.363420	208898.0
350	351	320	0.285081	150860.0
446	447	754	0.178394	206191.0
544	545	764	0.337769	135112.0
636	637	420	0.323984	140063.0
646	647	498	0.658087	128818.0
698	699	775	0.385100	190717.0
723	724	465	0.658173	140008.0
725	726	737	0.136048	205404.0
730	731	626	0.762245	209901.0
766	767	473	0.611750	195004.0
862	863	792	0.399555	208406.0
credit_inquiries_last_6_months credit_limit				
credit_score_range \				
1		2.0	1250.0	550-599
19		3.0	1500.0	600-649
25		2.0	60000.0	750-799
38		3.0	40000.0	700-749
93		2.0	40000.0	700-749
204		0.0	500.0	300-449
271		1.0	40000.0	700-749
301		4.0	40000.0	700-749
330		4.0	60000.0	750-799
350		0.0	500.0	300-449
446		2.0	60000.0	750-799
544		2.0	60000.0	750-799
636		4.0	500.0	300-449
646		3.0	750.0	450-499
698		2.0	60000.0	750-799
723		3.0	750.0	450-499

725	4.0	40000.0	700-749
730	2.0	1500.0	600-649
766	1.0	750.0	450-499
862	3.0	60000.0	750-799

	credit_limit_mode
1	1250.0
19	1500.0
25	60000.0
38	40000.0
93	40000.0
204	500.0
271	40000.0
301	40000.0
330	60000.0
350	500.0
446	60000.0
544	60000.0
636	500.0
646	750.0
698	60000.0
723	750.0
725	40000.0
730	1500.0
766	750.0
862	60000.0

```
df_cp_upd_3.loc[df_cp_upd_3['outstanding_debt'] >
df_cp_upd_3['credit_limit'], 'outstanding_debt']
```

1	161644.0
19	205014.0
25	190838.0
38	122758.0
93	137058.0
204	187849.0
271	154568.0
301	122402.0
330	208898.0
350	150860.0
446	206191.0
544	135112.0
636	140063.0
646	128818.0
698	190717.0

```
723    140008.0
725    205404.0
730    209901.0
766    195004.0
862    208406.0
```

```
Name: outstanding_debt, dtype: float64
```

```
df_cp_upd_3.loc[df_cp_upd_3['outstanding_debt'] >
df_cp_upd_3['credit_limit'], 'outstanding_debt'] =
df_cp_upd_3['credit_limit']
```

```
df_cp_upd_3.loc[[204, 544]]
```

	cust_id	credit_score	credit_utilisation	outstanding_debt \
204	205	303	0.364360	500.0
544	545	764	0.337769	60000.0

	credit_inquiries_last_6_months	credit_limit	credit_score_range \
204	0.0	500.0	300-449
544	2.0	60000.0	750-799

	credit_limit_mode
204	500.0
544	60000.0

```
df_cp_upd_3[df_cp_upd_3.outstanding_debt > df_cp_upd_3.credit_limit]
```

```
Empty DataFrame
```

```
Columns: [cust_id, credit_score, credit_utilisation, outstanding_debt,
credit_inquiries_last_6_months, credit_limit, credit_score_range,
credit_limit_mode]
Index: []
```

```
df_cp_upd_3.describe()
```

	cust_id	credit_score	credit_utilisation	outstanding_debt
count	1000.000000	1000.000000	1000.000000	1000.000000
mean	500.500000	589.182000	0.498950	6850.084000
std	288.819436	152.284929	0.233139	10683.473561
min	1.000000	300.000000	0.103761	33.000000
25%	250.750000	460.000000	0.293917	221.000000
50%	500.500000	601.500000	0.487422	541.500000

75%	750.250000	738.000000	0.697829	10924.500000
max	1000.000000	799.000000	0.899648	60000.000000
	credit_inquiries_last_6_months	credit_limit	credit_limit_mode	
count	1000.000000	1000.000000	1000.000000	
mean	1.955000	19733.75000	19912.500000	
std	1.414559	24717.43818	24840.914633	
min	0.000000	500.00000	500.000000	
25%	1.000000	750.00000	750.000000	
50%	2.000000	1500.00000	1500.000000	
75%	3.000000	40000.00000	40000.000000	
max	4.000000	60000.00000	60000.000000	

4 - Data Exploration : Visualizing Correlation in Credit Score Table

```
df_cust.head(2)
```

	cust_id	name	gender	age	location	occupation	\
0	1	Manya Acharya	Female	51.0	City	Business Owner	
1	2	Anjali Pandey	Female	47.0	City	Consultant	

	annual_income	marital_status	age_group
0	358211.0	Married	49-65
1	65172.0	Single	26-48

```
df_cp_upd_3.head(2)
```

	cust_id	credit_score	credit_utilisation	outstanding_debt	\
0	1	749	0.585171	19571.0	
1	2	587	0.107928	1250.0	

	credit_inquiries_last_6_months	credit_limit	credit_score_range	\
0	0.0	40000.0	700-749	
1	2.0	1250.0	550-599	

	credit_limit_mode
0	40000.0
1	1250.0

```
df_merged = df_cust.merge(df_cp_upd_3, on='cust_id', how='inner')
df_merged.head(2)
```

	cust_id	name	gender	age	location	occupation	\
0	1	Manya Acharya	Female	51.0	City	Business Owner	
1	2	Anjali Pandey	Female	47.0	City	Consultant	

	annual_income	marital_status	age_group	credit_score
0	358211.0	Married	49-65	749
0.585171				
1	65172.0	Single	26-48	587
0.107928				

	outstanding_debt	credit_inquiries_last_6_months	credit_limit	\
0	19571.0	0.0	40000.0	
1	1250.0	2.0	1250.0	

	credit_score_range	credit_limit_mode
0	700-749	40000.0
1	550-599	1250.0

```
numerical_cols = ['credit_score', 'credit_utilisation',
'outstanding_debt', 'credit_limit', 'annual_income', 'age']
```

```
corr_matx = df_merged[numerical_cols].corr()
corr_matx
```

	credit_score	credit_utilisation	outstanding_debt
\			
credit_score	1.000000	-0.070445	0.680654
credit_utilisation	-0.070445	1.000000	0.192838
outstanding_debt	0.680654	0.192838	1.000000
credit_limit	0.847952	-0.080493	0.810581
annual_income	0.575685	-0.086816	0.555077
age	0.444917	-0.027713	0.444301

	credit_limit	annual_income	age
credit_score	0.847952	0.575685	0.444917
credit_utilisation	-0.080493	-0.086816	-0.027713
outstanding_debt	0.810581	0.555077	0.444301
credit_limit	1.000000	0.684627	0.510993
annual_income	0.684627	1.000000	0.618136
age	0.510993	0.618136	1.000000


Creating a list of numerical columns you're interested in — basically telling Python, "These are the columns I want to study for relationships."

The above code does 3 things:

- `df_merged[numerical_cols]` — selects just those columns from the DataFrame `df_merged`.
- `.corr()` — calculates the correlation between every pair of those columns.
- Stores the result in a new variable called `corr_matx`.

```
# Creating a heatmap of the correlation matrix
```

```
plt.figure(figsize=(5, 3))  
sns.heatmap(corr_matx, annot=True, fmt=".2f", cmap='coolwarm',  
linewidths=0.8)  
plt.title('Correlation Plot')  
plt.show()
```



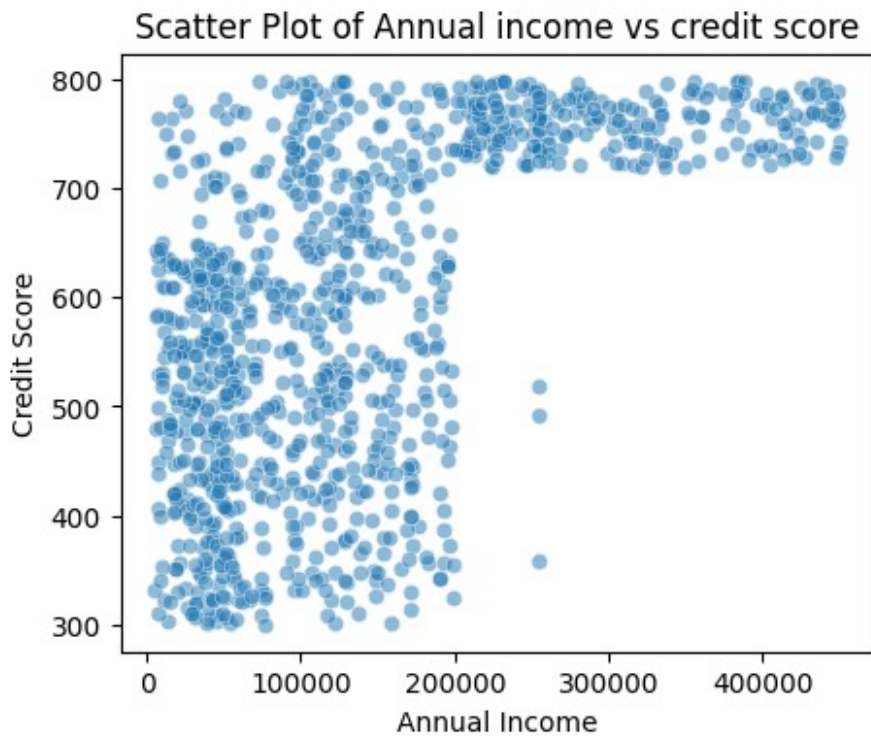
You can see a **high correlation** between credit limit & credit score (~0.85), credit limit & annual income.

This correlation table can be used for further analysis. It shows if one variable has relationship with the other variable

```
# Checking if there is any relation between annual_income and credit  
score
```



```
plt.figure(figsize=(5, 4))
sns.scatterplot(x='annual_income', y='credit_score', data=df_merged,
alpha=0.5)
plt.title('Scatter Plot of Annual income vs credit score')
plt.xlabel('Annual Income')
plt.ylabel('Credit Score')
plt.show()
```



No clear pattern observed.

Transactions Table

```
df_trans.head(10)
```

	tran_id	cust_id	tran_date	tran_amount	platform \
0	1	705	2023-01-01	63	Flipkart
1	2	385	2023-01-01	99	Alibaba
2	3	924	2023-01-01	471	Shopify
3	4	797	2023-01-01	33	Shopify
4	5	482	2023-01-01	68	Amazon
5	6	527	2023-01-01	38	Shopify
6	7	388	2023-01-01	720	Alibaba
7	8	8	2023-01-01	140	Shopify
8	9	939	2023-01-01	144	Alibaba

9	10	228	2023-01-01	836	Ebay
---	----	-----	------------	-----	------

	product_category	payment_type
0	Electronics	Phonepe
1	Fashion & Apparel	Credit Card
2	Sports	Phonepe
3	Fashion & Apparel	Gpay
4	Fashion & Apparel	Net Banking
5	Fashion & Apparel	Debit Card
6	Electronics	Credit Card
7	Kitchen Appliances	Gpay
8	Beauty & Personal Care	Phonepe
9	Electronics	Gpay

```
df_trans.shape
```

```
(500000, 7)
```

```
df_trans.isnull().sum()
```

```
tran_id          0
cust_id          0
tran_date        0
tran_amount      0
platform        4941
product_category  0
payment_type     0
dtype: int64
```

1 - Handling Null Values

```
df_trans[df_trans.platform.isnull()]
```

	tran_id	cust_id	tran_date	tran_amount	platform	product_category
355	356	58	2023-01-01	237	None	Electronics
418	419	383	2023-01-01	338	None	Electronics
607	608	421	2023-01-01	700	None	Electronics
844	845	945	2023-01-01	493	None	Sports
912	913	384	2023-01-01	85	None	Fashion & Apparel
...
499579	499580	924	2023-09-05	31	None	Fashion &

```

Apparel
499646  499647      944  2023-09-05      58445      None  Fashion &
Apparel
499725  499726      620  2023-09-05         15      None
Sports
499833  499834      616  2023-09-05         97      None  Fashion &
Apparel
499997  499998        57  2023-09-05        224      None  Garden &
Outdoor

      payment_type
355      Net Banking
418      Credit Card
607          Phonepe
844      Credit Card
912          Phonepe
...          ...
499579          Gpay
499646          Phonepe
499725  Net Banking
499833  Credit Card
499997          Phonepe

[4941 rows x 7 columns]

```

For above, we cannot replace null values with Mean or Median.

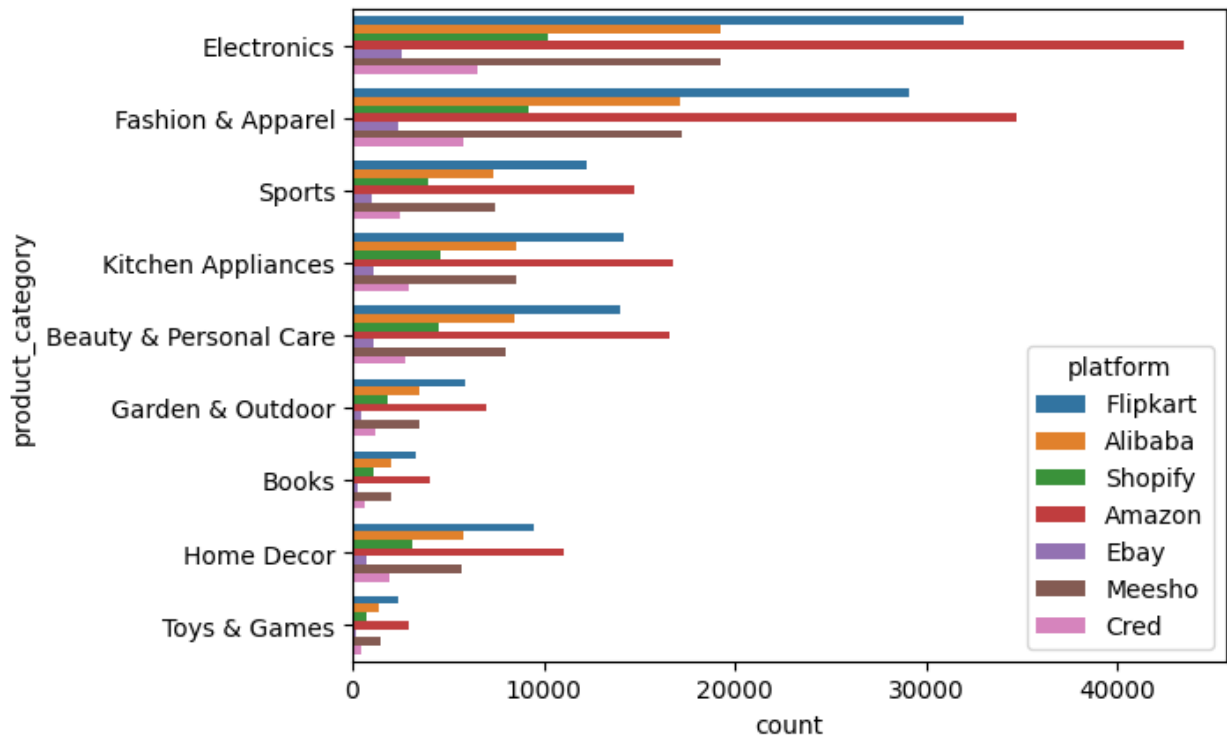
We'll go with **Mode**. And by segmenting the **product_category** we can assign the platform respectively.

```

sns.countplot(y = 'product_category', hue = 'platform', data =
df_trans)

<Axes: xlabel='count', ylabel='product_category'>

```



In the above chart, we can see that in all *product_category*, **Amazon** is the platform that is used the most for making purchases.

For handling null values in **platform**, we can just replace them using **Amazon** as a product platform just because it is used most frequently.

```
df_trans.platform.mode()[0]
'Amazon'
df_trans['platform'].fillna(df_trans.platform.mode()[0], inplace=True)
df_trans.isnull().sum()
tran_id      0
cust_id      0
tran_date    0
tran_amount  0
platform     0
product_category  0
payment_type  0
dtype: int64
```

2 - Treating Outliers for tran_amount

```
df_trans.describe()
```

	tran_id	cust_id	tran_amount
count	500000.000000	500000.000000	500000.000000
mean	250000.500000	501.400428	3225.20733
std	144337.711635	288.641924	13098.74276
min	1.000000	1.000000	0.000000
25%	125000.750000	252.000000	64.000000
50%	250000.500000	502.000000	141.000000
75%	375000.250000	752.000000	397.000000
max	500000.000000	1000.000000	69999.000000

We can see transactions with **0 amount**. These seem to be invalid

```
df_trans_zero = df_trans[df_trans.tran_amount == 0]  
df_trans_zero.head(3)
```

	tran_id	cust_id	tran_date	tran_amount	platform
product_category \					
120	121	440	2023-01-01	0	Amazon
Electronics					
141	142	839	2023-01-01	0	Amazon
Electronics					
517	518	147	2023-01-01	0	Amazon
Electronics					

	payment_type
120	Credit Card
141	Credit Card
517	Credit Card

```
df_trans_zero.shape
```

```
(4734, 7)
```

```
df_trans_zero.platform.value_counts()
```

```
platform  
Amazon    4734  
Name: count, dtype: int64
```

```
df_trans_zero.product_category.value_counts()
```

```
product_category  
Electronics    4734  
Name: count, dtype: int64
```

```
df_trans_zero.payment_type.value_counts()
```

```
payment_type
Credit Card    4734
Name: count, dtype: int64
```

It appears that when **platform=Amazon, product_category=Eletronics and payment_type=Credit Card**, at that time we get all these zero transactions.

We need to find other transactions in this group and find its median to replace these zero values. We are not using mean because we can see some outliers as well in this column.

```
df_trans_1 =
df_trans[(df_trans.platform=='Amazon')&(df_trans.product_category=="Electronics")&(df_trans.payment_type=="Credit Card")]
df_trans_1.shape
```

```
(15637, 7)
```

```
df_trans_1[df_trans_1.tran_amount>0]
```

	tran_id	cust_id	tran_date	tran_amount	platform
product_category \					
109	110	887	2023-01-01	635	Amazon
Electronics					
173	174	676	2023-01-01	60439	Amazon
Electronics					
190	191	763	2023-01-01	697	Amazon
Electronics					
263	264	528	2023-01-01	421	Amazon
Electronics					
311	312	936	2023-01-01	537	Amazon
Electronics					
...
...					
499766	499767	723	2023-09-05	909	Amazon
Electronics					
499793	499794	586	2023-09-05	304	Amazon
Electronics					
499812	499813	688	2023-09-05	425	Amazon
Electronics					
499860	499861	373	2023-09-05	480	Amazon
Electronics					
499885	499886	520	2023-09-05	643	Amazon
Electronics					

```
payment_type
109    Credit Card
```

```

173      Credit Card
190      Credit Card
263      Credit Card
311      Credit Card
...
499766   Credit Card
499793   Credit Card
499812   Credit Card
499860   Credit Card
499885   Credit Card

[10903 rows x 7 columns]

median_to_replace =
df_trans_1[df_trans_1.tran_amount>0].tran_amount.median()
median_to_replace

554.0

df_trans['tran_amount'].replace(0,median_to_replace, inplace=True)

df_trans[df_trans.tran_amount==0]

Empty DataFrame
Columns: [tran_id, cust_id, tran_date, tran_amount, platform,
product_category, payment_type]
Index: []

```

No 0 values is present in **tran_amount** column

```

df_trans.tran_amount.describe()

count      500000.000000
mean        3230.452602
std         13097.561071
min           2.000000
25%          66.000000
50%         146.000000
75%         413.000000
max        69999.000000
Name: tran_amount, dtype: float64

df_trans[df_trans['tran_amount']<1000].describe()


```

	tran_id	cust_id	tran_amount
count	475000.000000	475000.000000	475000.000000
mean	250041.699922	501.375499	240.667608
std	144285.259913	288.606185	244.487110
min	1.000000	1.000000	2.000000
25%	125126.750000	252.000000	63.000000
50%	250100.500000	502.000000	131.000000

75%	374928.250000	751.000000	348.000000
max	500000.000000	1000.000000	999.000000

```
Q1, Q3 = df_trans['tran_amount'].quantile([0.25, 0.75])
IQR = Q3 - Q1
lower = Q1 - 2 * IQR      # 2 instead of 1.5 (little flexible for business)
upper = Q3 + 2 * IQR
```

```
lower, upper
```

```
(-628.0, 1107.0)
```

```
df_trans[df_trans.tran_amount<upper].tran_amount.max()
```

```
999
```

```
df_trans[df_trans.tran_amount<upper].tran_amount.min()
```

```
2
```

```
df_trans_outliers = df_trans[df_trans.tran_amount>=upper]
df_trans_outliers
```

	tran_id	cust_id	tran_date	tran_amount	platform	\
26	27	380	2023-01-01	61963	Shopify	
49	50	287	2023-01-01	57869	Amazon	
94	95	770	2023-01-01	52881	Ebay	
104	105	549	2023-01-01	58574	Flipkart	
113	114	790	2023-01-01	51669	Shopify	
...	
499742	499743	868	2023-09-05	55131	Meesho	
499888	499889	614	2023-09-05	59679	Meesho	
499900	499901	811	2023-09-05	60184	Flipkart	
499966	499967	662	2023-09-05	54678	Meesho	
499996	499997	569	2023-09-05	53022	Meesho	

	product_category	payment_type
26	Beauty & Personal Care	Credit Card
49	Toys & Games	Gpay
94	Kitchen Appliances	Credit Card
104	Fashion & Apparel	Gpay
113	Kitchen Appliances	Credit Card
...
499742	Fashion & Apparel	Gpay
499888	Fashion & Apparel	Net Banking
499900	Sports	Debit Card
499966	Sports	Gpay
499996	Fashion & Apparel	Net Banking

```
[25000 rows x 7 columns]
```



```
df_trans_normal = df_trans[df_trans.tran_amount < upper]
df_trans_normal
```

	tran_id	cust_id	tran_date	tran_amount	platform	\
0	1	705	2023-01-01	63	Flipkart	
1	2	385	2023-01-01	99	Alibaba	
2	3	924	2023-01-01	471	Shopify	
3	4	797	2023-01-01	33	Shopify	
4	5	482	2023-01-01	68	Amazon	
...
499994	499995	679	2023-09-05	59	Ebay	
499995	499996	791	2023-09-05	43	Amazon	
499997	499998	57	2023-09-05	224	Amazon	
499998	499999	629	2023-09-05	538	Flipkart	
499999	500000	392	2023-09-05	346	Amazon	

	product_category	payment_type
0	Electronics	Phonepe
1	Fashion & Apparel	Credit Card
2	Sports	Phonepe
3	Fashion & Apparel	Gpay
4	Fashion & Apparel	Net Banking
...
499994	Beauty & Personal Care	Gpay
499995	Books	Phonepe
499997	Garden & Outdoor	Phonepe
499998	Home Decor	Gpay
499999	Kitchen Appliances	Net Banking

[475000 rows x 7 columns]

```
tran_mean_per_category = df_trans_normal.groupby("product_category")
["tran_amount"].mean()
tran_mean_per_category
```

product_category	
Beauty & Personal Care	92.167205
Books	29.553515
Electronics	510.172685
Fashion & Apparel	64.553463
Garden & Outdoor	125.630277
Home Decor	302.487561
Kitchen Appliances	176.773288
Sports	269.181631
Toys & Games	50.333298

Name: tran_amount, dtype: float64

```
df_trans.loc[df_trans_outliers.index]
```

	tran_id	cust_id	tran_date	tran_amount	platform	\
26	27	380	2023-01-01	61963	Shopify	

49	50	287	2023-01-01	57869	Amazon
94	95	770	2023-01-01	52881	Ebay
104	105	549	2023-01-01	58574	Flipkart
113	114	790	2023-01-01	51669	Shopify
...
499742	499743	868	2023-09-05	55131	Meesho
499888	499889	614	2023-09-05	59679	Meesho
499900	499901	811	2023-09-05	60184	Flipkart
499966	499967	662	2023-09-05	54678	Meesho
499996	499997	569	2023-09-05	53022	Meesho

	product_category	payment_type
26	Beauty & Personal Care	Credit Card
49	Toys & Games	Gpay
94	Kitchen Appliances	Credit Card
104	Fashion & Apparel	Gpay
113	Kitchen Appliances	Credit Card
...
499742	Fashion & Apparel	Gpay
499888	Fashion & Apparel	Net Banking
499900	Sports	Debit Card
499966	Sports	Gpay
499996	Fashion & Apparel	Net Banking

[25000 rows x 7 columns]

```
df_trans.loc[df_trans_outliers.index, 'tran_amount'] =
df_trans_outliers['product_category'].map(tran_mean_per_category)

df_trans.loc[df_trans_outliers.index]
```

	tran_id	cust_id	tran_date	tran_amount	platform \
26	27	380	2023-01-01	92.167205	Shopify
49	50	287	2023-01-01	50.333298	Amazon
94	95	770	2023-01-01	176.773288	Ebay
104	105	549	2023-01-01	64.553463	Flipkart
113	114	790	2023-01-01	176.773288	Shopify
...
499742	499743	868	2023-09-05	64.553463	Meesho
499888	499889	614	2023-09-05	64.553463	Meesho
499900	499901	811	2023-09-05	269.181631	Flipkart
499966	499967	662	2023-09-05	269.181631	Meesho
499996	499997	569	2023-09-05	64.553463	Meesho

	product_category	payment_type
26	Beauty & Personal Care	Credit Card
49	Toys & Games	Gpay
94	Kitchen Appliances	Credit Card
104	Fashion & Apparel	Gpay
113	Kitchen Appliances	Credit Card

```

...
499742      Fashion & Apparel      ...      Gpay
499888      Fashion & Apparel      Net Banking
499900              Sports      Debit Card
499966              Sports      Gpay
499996      Fashion & Apparel      Net Banking

```


[25000 rows x 7 columns]

We now got rid of **outliers** from **tran_amount** column.

```

sns.histplot(x='tran_amount', data=df_trans, bins=20, kde=True)
<Axes: xlabel='tran_amount', ylabel='Count'>

```



Above shows the histogram of transactions after the removal of outliers.

We can see that distribution is right skewed. Transaction amount now is **< 1000**

Data Visualization : Payment Type Distribution

```
df_trans.head(3)
```

```


    tran_id  cust_id  tran_date  tran_amount  platform
product_category \
0          1      705  2023-01-01         63.0  Flipkart
Electronics
1          2      385  2023-01-01         99.0  Alibaba Fashion &
Apparel
2          3      924  2023-01-01        471.0  Shopify
Sports

    payment_type
0      Phonepe
1  Credit Card
2      Phonepe

sns.countplot(x=df_trans.payment_type, stat='percent')

<Axes: xlabel='payment_type', ylabel='percent'>

```



The above plot shows the **Distribution of payment types across age groups**.

```

df_merged_2 = df_merged.merge(df_trans, on='cust_id', how='inner')
df_merged_2.head(3)

```

	cust_id	name	gender	age	location	occupation	\
0	1	Manya Acharya	Female	51.0	City	Business Owner	
1	1	Manya Acharya	Female	51.0	City	Business Owner	
2	1	Manya Acharya	Female	51.0	City	Business Owner	

	annual_income	marital_status	age_group	credit_score	...	\
0	358211.0	Married	49-65	749	...	
1	358211.0	Married	49-65	749	...	
2	358211.0	Married	49-65	749	...	

	credit_inquiries_last_6_months	credit_limit	credit_score_range	\
0	0.0	40000.0	700-749	
1	0.0	40000.0	700-749	
2	0.0	40000.0	700-749	

	credit_limit_mode	tran_id	tran_date	tran_amount	platform	\
0	40000.0	1283	2023-01-01	30.0	Shopify	
1	40000.0	1382	2023-01-01	96.0	Amazon	
2	40000.0	1521	2023-01-01	86.0	Meesho	

	product_category	payment_type
0	Fashion & Apparel	Net Banking
1	Sports	Debit Card
2	Garden & Outdoor	Gpay


[3 rows x 22 columns]

df_merged_2.shape

(500000, 22)

```
plt.figure(figsize=(5, 4))
sns.countplot(x='age_group', hue='payment_type', data=df_merged_2,
palette='Set3')
plt.title('Distribution of Payment-Types across Age-Groups')
plt.xlabel('Age Group')
plt.ylabel('Count')
plt.xticks(rotation=0)
plt.legend(title='Payment Type', loc='upper right')

plt.show()
```




From above analysis, we can see that **age group 18-25** has less exposure to credit cards compared to other groups

```
fig, (ax1, ax2) = plt.subplots(1,2, figsize=(12,5))

sns.countplot(x='age_group', hue="product_category", data=df_merged_2,
ax=ax1)
ax1.set_title("Product-Category Count By Age-Group")
ax1.set_xlabel("Age Group")
ax1.set_ylabel("Count")
ax1.legend(title="Product Category", loc='upper right')

sns.countplot(x='age_group', hue="platform", data=df_merged_2, ax=ax2)
ax2.set_title("Platform Count By Age-Group")
ax2.set_xlabel("Age Group")
ax2.set_ylabel("Count")
ax2.legend(title="Product Category", loc='upper right')

plt.show()
```



Observations:

- **Top 3** purchasing categories of customers in age group (18 -25) : Electronics, Fashion & Apparel, Beauty & personal care
- **Top 3** platforms : Amazon, Flipkart, Alibaba

Data Visualization : Average Transaction Amount

```
# List of categorical columns
cat_cols = ['payment_type', 'platform', 'product_category',
            'marital_status', 'age_group']

num_rows = 3
# Create subplots
fig, axes = plt.subplots(num_rows, 2, figsize=(12, 4 * num_rows))

# Flatten the axes array to make it easier to iterate
axes = axes.flatten()

# Create subplots for each categorical column
for i, cat_col in enumerate(cat_cols):
    # Calculate the average annual income for each category
    avg_tran_amount_by_category = df_merged_2.groupby(cat_col)
    ['tran_amount'].mean().reset_index()

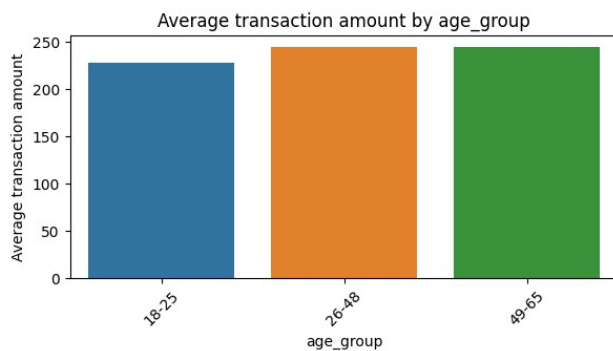
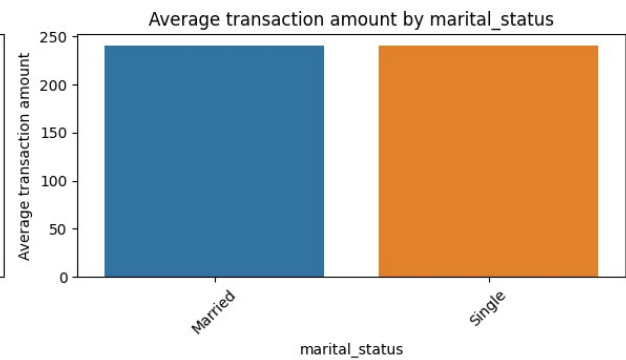
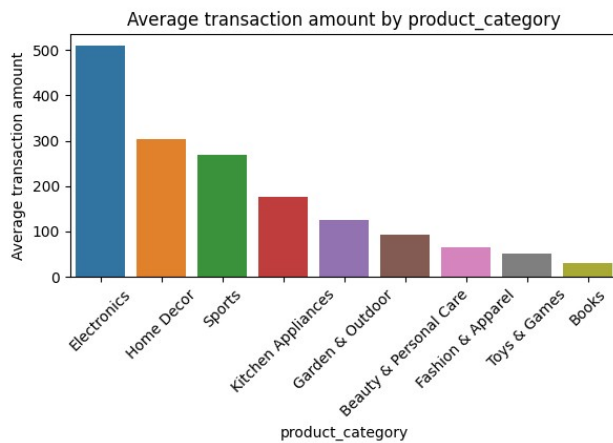


    # Sort the data by 'annual_income' before plotting
    sorted_data =
    avg_tran_amount_by_category.sort_values(by='tran_amount',
    ascending=False)

    sns.barplot(x=cat_col, y='tran_amount', data=sorted_data, ci=None,
```

```
ax=axes[i], palette='tab10')
    axes[i].set_title(f'Average transaction amount by {cat_col}')
    axes[i].set_xlabel(cat_col)
    axes[i].set_ylabel('Average transaction amount')

    # Rotate x-axis labels for better readability
    axes[i].set_xticklabels(axes[i].get_xticklabels(), rotation=45)

# Hide any unused subplots
for i in range(len(cat_cols), len(axes)):
    fig.delaxes(axes[i])
plt.tight_layout()
plt.show()
```

Let us do further analysis on age group to figure out their average income, credit limit, credit score etc

```
# Group the data by age group and calculate the average credit_limit and credit_score
```

```
age_group_metrics = df_merged.groupby('age_group')[['annual_income', 'credit_limit', 'credit_score']].mean().reset_index()
age_group_metrics
```

	age_group	annual_income	credit_limit	credit_score
0	18-25	36969.670732	1130.081301	484.451220
1	26-48	145437.104938	20560.846561	597.569665
2	49-65	259786.192513	41699.197861	701.524064

```
# Create subplots
```

```
fig, (ax1, ax2, ax3) = plt.subplots(1, 3, figsize=(12, 4))
```

```
# Plot 1: Average annual income by age group
```

```
sns.barplot(x='age_group', y='annual_income', data=age_group_metrics,
palette='tab10', ax=ax1)
ax1.set_title('Average Annual Income by Age Group')
ax1.set_xlabel('Age Group')
ax1.set_ylabel('Average Annual Income')
ax1.tick_params(axis='x', rotation=0)
```


```
# Plot 2: Average Max Credit Limit by Age Group
```

```
sns.barplot(x='age_group', y='credit_limit', data=age_group_metrics,
palette='hls', ax=ax2)
ax2.set_title('Average Credit Limit by Age Group')
ax2.set_xlabel('Age Group')
ax2.set_ylabel('Average Credit Limit')
ax2.tick_params(axis='x', rotation=0)
```

```
# Plot 3: Average Credit Score by Age Group
```

```
sns.barplot(x='age_group', y='credit_score', data=age_group_metrics,
palette='viridis', ax=ax3)
ax3.set_title('Average Credit Score by Age Group')
ax3.set_xlabel('Age Group')
ax3.set_ylabel('Average Credit Score')
ax3.tick_params(axis='x', rotation=0)
```

```
plt.tight_layout()
plt.show()
```



Finalize Target Market For a Trial Credit Card Launch

- People with **age group of 18-25** accounts to ~26% of customer base in the data
- Avg annual income of this group is **<50k**
- They **don't have much credit history** which is getting reflected in their credit score and credit limit
- **Usage of credit cards** as payment type is relatively **low** compared to other groups
- **Top 3** most shopping products categories : Electronics, Fashion & Apparel, Beauty & Personal care